

Velkommen til

Network Management using Mostly Open Source

Henrik Lund Kramshøj
hlik@solidonetworks.com

<http://www.solidonetworks.com>

Slides are available as PDF

*Network and Internet is an integral part of our everyday lives,
but how does one ensure that the network works perfectly.*

Introduce essential tools for network management

Prove that open source is critical for network management

Present resources for others to follow

Expect you to be administrators of IP networks, in some way

Presentation is based on the experience from an ISP viewpoint

Walk through of the essential tools and skills you need to acquire


Contents


The problem: what is network management

Components of a solution

The solution embrace, extend and proliferate :-)

The problem: what is network management





Core routers at Interxion in Ballerup

Second major site in Luxembourg


Internet connections multiple 10Gbit at each site

New servers - every week, new switches - every month

Support systems, APC ATS, APC PDU, routers, switches, jump hosts, monitoring, logging servers, support system, ...

Scheduled down time or outages is not an option anymore

AS12617 - high level BGP



Source: <http://bgp.he.net/AS12617>

Step 1: configure devices properly

You should always configure your devices properly

Turn on SNMP, probably SNMPv2

Turn on LLDP Link Layer Discovery Protocol,
like CDP but vendor-neutral

http://en.wikipedia.org/wiki/Link_Layer_Discovery_Protocol

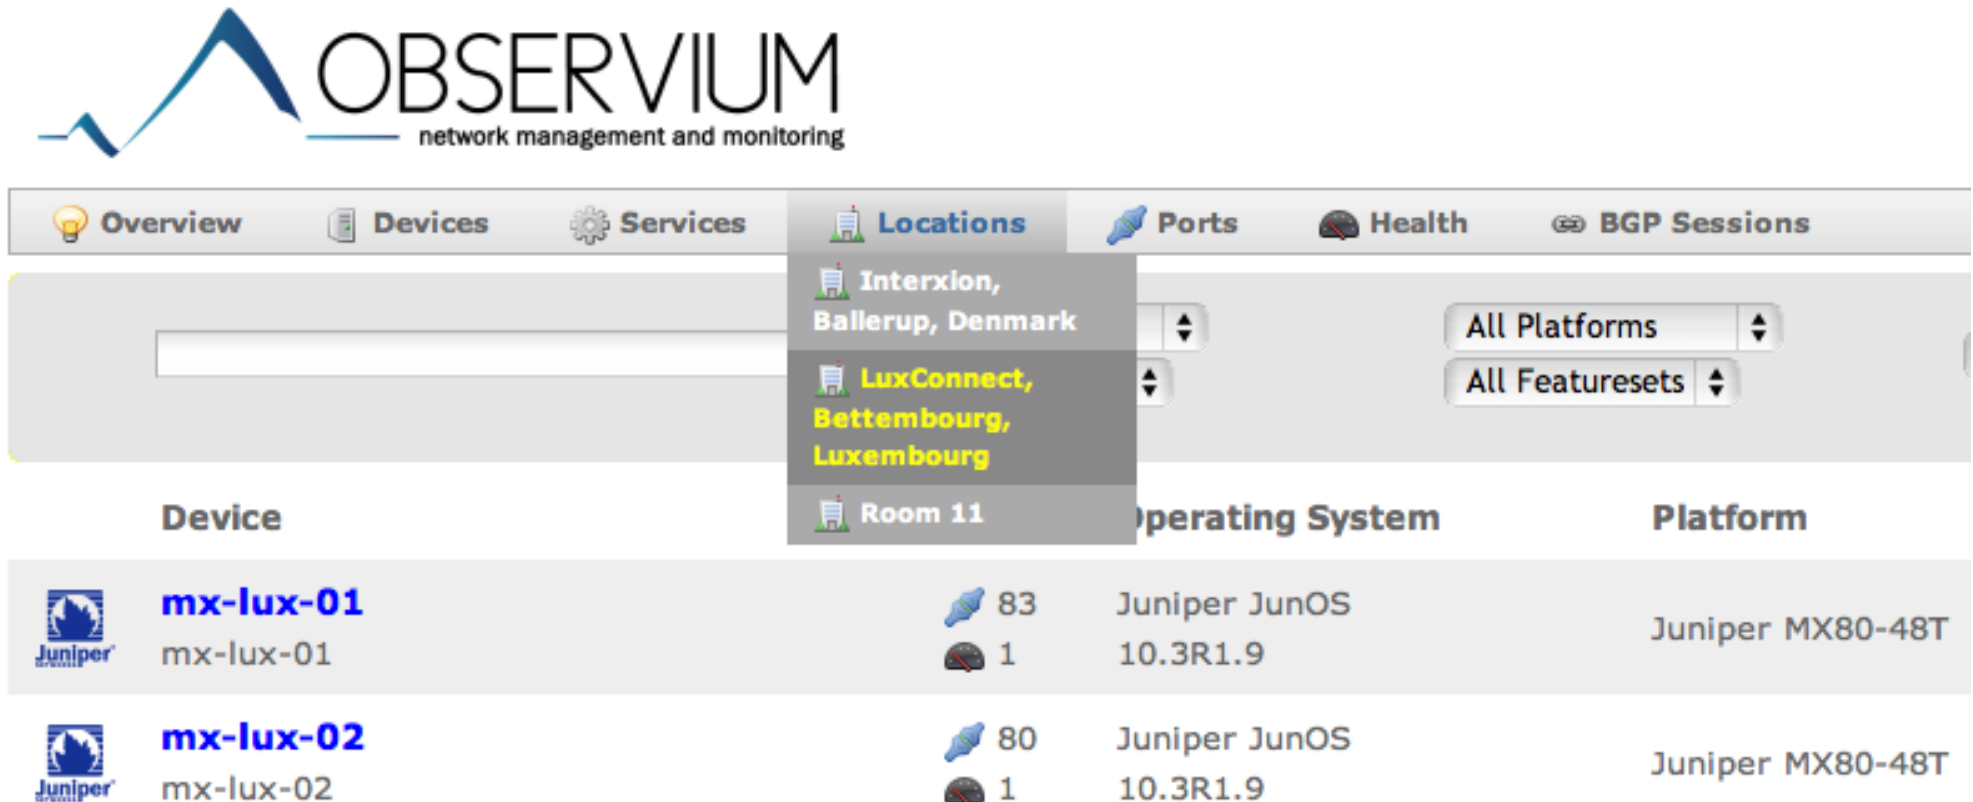
Syslog - you know this, nuff said

And updated firmware, HTTPS and SSH only etc. the usual stuff







Config example: SNMP

```
snmp {  
    description "Solido Networks SRX-CPH-02";  
    location "Interxion, Ballerup, Denmark";  
    contact "noc@solido.net";  
    community yourcommunitynotmine {  
        authorization read-only;  
        clients {  
            10.1.1.1/32;  
            10.1.2.2/32;  
        }  
    }  
}
```

Location, location, location

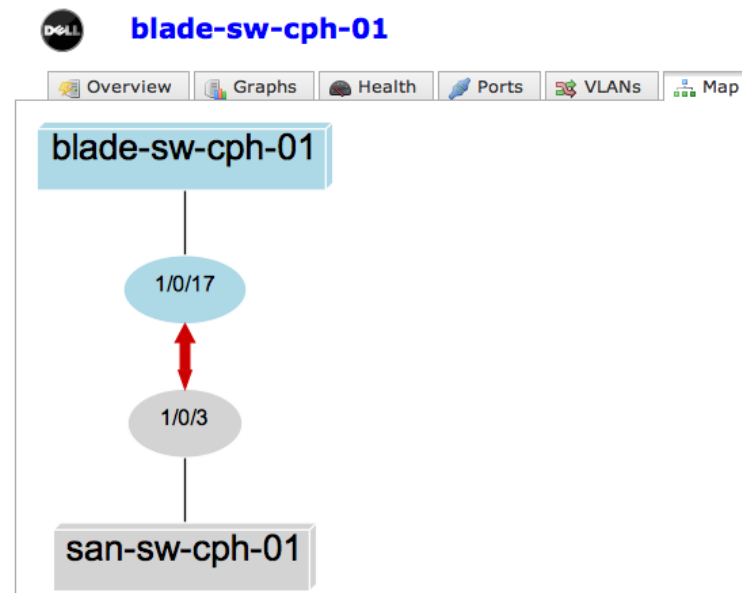


The screenshot shows the Observium web interface. At the top is the Observium logo and the text "network management and monitoring". Below this is a navigation bar with tabs: Overview, Devices, Services, Locations, Ports, Health, and BGP Sessions. The "Locations" tab is active, and its dropdown menu is open, showing three options: "Interxion, Ballerup, Denmark", "LuxConnect, Bettembourg, Luxembourg" (highlighted in yellow), and "Room 11". To the right of the dropdown are two filters: "All Platforms" and "All Featuresets". Below the navigation bar is a table with the following columns: Device, Operating System, and Platform. The table contains two rows of data for Juniper devices.

Device	Operating System	Platform
 mx-lux-01 mx-lux-01	 83  1 Juniper JunOS 10.3R1.9	Juniper MX80-48T
 mx-lux-02 mx-lux-02	 80  1 Juniper JunOS 10.3R1.9	Juniper MX80-48T

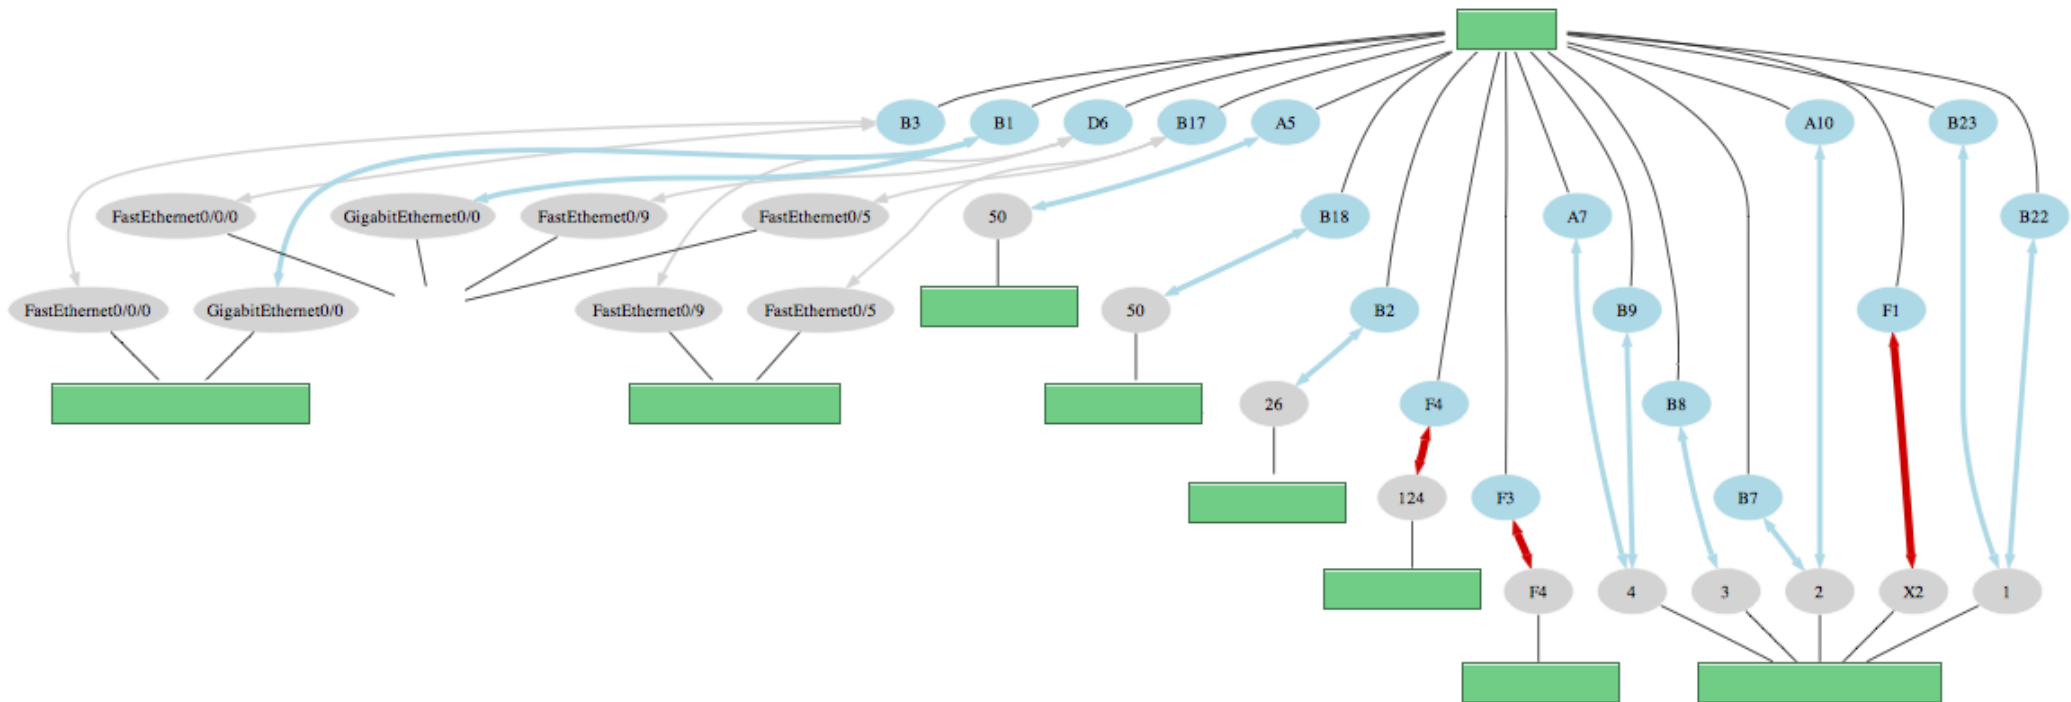
Observium picks up the location from SNMP :-)

Dell 8024F switch LLDP



```
interface ethernet 1/xg17
mtu 9216
lldp transmit-tlv port-desc sys-name sys-desc sys-cap
lldp transmit-mgmt
exit
```

LLDP spaghetti?



LLDP is needed!

LLDP trick using tcpdump

```
[hlk@ljlh ~]$ sudo tcpdump -i eth0 ether proto 0x88cc
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on eth0, link-type EN10MB (Ethernet), capture size 96 bytes
.... wait for it ....
11:03:55.395064 00:1c:23:80:49:ff (oui Unknown) > 01:80:c2:00:00:0e (oui Unknown),
ethertype Unknown (0x88cc), length 60:
0x0000:  0207 0400 1c23 8049 fd04 0705 312f 302f  ....#.I....1/0/
0x0010:  3331 0602 0078 0000 0000 0000 0000 0000  31...x.....
0x0020:  0000 0000 0000 0000 0000 0000 0000 0000  .....
```

1 packets captured
3 packets received by filter
0 packets dropped by kernel

I know **for sure** that this server is in Unit 1 port 31!

The internet is a collaborative experiment, hippies connecting their routers and liberally exchanging information - sometime money is exchanged also :-)

Main points, operation, administration, maintenance

Configure your network equipment, provisioning

Monitor your network

Control your network - react to events

Troubleshoot your network

Before running in production, and when troubleshooting

Ping and traceroute - you know these, Unix traceroute ICMP/UDP

Nping, Nmap, Mtr, TCP traceroute, hping, icmpush ...

Download Backtrack Linux now, it is your network toolbox

Huge number of goodies on Backtrack for network management!

<http://backtrack-linux.org/>

Learn Unix - yes, Linux/Unix is needed when working with networks

You need skills in sed/awk, cut, **expect**, grep, sort, Perl/Python/Ruby at least one scripting language

Conserver is an application that allows multiple users to watch a serial console at the same time. It can log the data, allows users to take write-access of a console (one at a time), and has a variety of bells and whistles to accentuate that basic functionality.

Watch the console!

A network device rebooted - what happened?

I accidentally the whole network, what now?

Serial consoles are not dead, and still very useful

`http://www.conserver.com/`



Soekris, 4-port serial card EUR59 / 430DKK + OpenBSD + conserver

```
### set the defaults for all the consoles
# these get applied before anything else
default * {
    # The '&' character is substituted with the console name
    logfile /var/consoles/&;
    # timestamps every hour with activity and break logging
    timestamp lhab;
    # include the 'full' default
    include full;
    # master server is localhost
    master localhost;
}
...
console portS1 {
    type device;
    device /dev/cua02; parity none; baud 57600;
    idlestring "#";
    idletimeout 5m;           # send a '#' every 5 minutes of idle
    timestamp "";            # no timestamps on this console
}
```

You will actually be able to say what happened at that device

MRTG The Multi Router Traffic Grapher - simple, great, fast

<http://oss.oetiker.ch/mrtg/>

Smokeping Network Latency measurements - network quality


<http://oss.oetiker.ch/smokeping/>

NFsen Netflow monitoring - turn on at selected routers/switches


Manual tools, My Traceroute, Nping

Routers in Luxembourg

Traffic Analysis for xe-0/0/3 -- mx-lux-01 Global Crossing




Traffic Analysis for xe-0/0/1 -- mx-lux-01 link to MX2




Run configmaker, indexmaker - almost done

Smokeping packet loss



Smokeping latency changed



Netflow is getting more important, more data share the same links

Accounting is important

Detecting DoS/DDoS and problems is essential

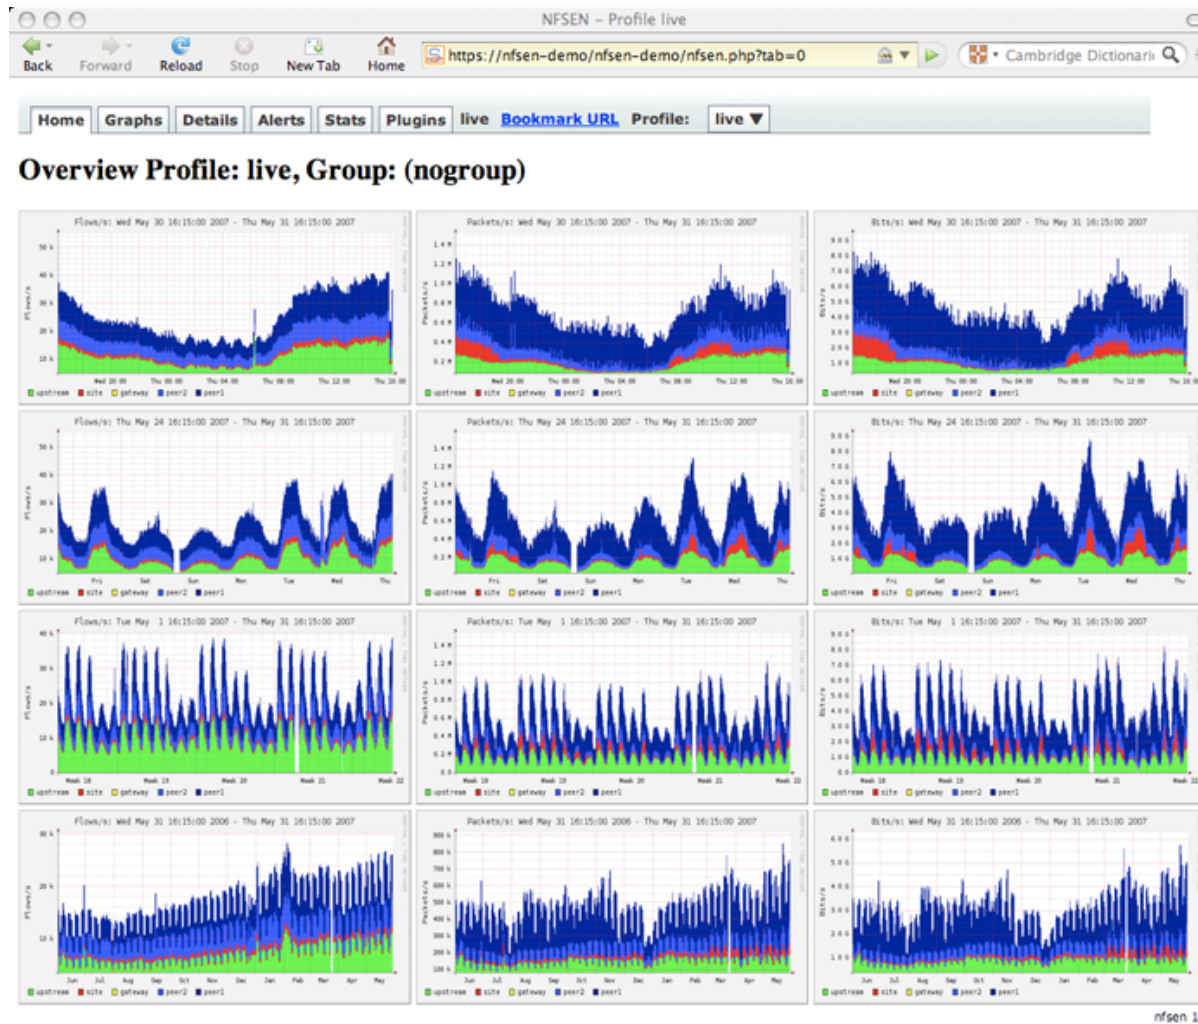
Netflow sampling is vital information - 123Mbit, but what kind of traffic

We use mostly NfSen, but are looking at various software packages

<http://nfsen.sourceforge.net/>

Currently also investigating sFlow - hopefully more fine grained

Netflow using NFSen



Netflow processing from the web interface

NFSN - Profile live May 31 2007 - 04:40

Back Forward Reload Stop New Tab Home <https://nfsen-demo/nfsen-demo/nfsen.php#processing> Cambridge Dictionarie

	peer2	gateway	site	upstream
3.3 k/s	76.2 k/s	66.9 k/s	7.0 k/s	621.0 /s
1.7 k/s	484.6 Mb/s	459.9 Mb/s	12.5 Mb/s	437.3 kb/s
11.7 Mb/s	1.0 /s	651.0 /s	600.8 /s	46.6 /s
0 /s	3.7 /s	6.2 Mb/s	6.1 Mb/s	36.4 kb/s
0 b/s	4.4 kb/s	467.1 /s	8.9 k/s	6.1 k/s
2.0 k/s	181.7 /s	613.3 /s	38.8 Mb/s	28.3 Mb/s
7.4 Mb/s	104.0 kb/s	2.9 Mb/s	6.4 k/s	94.2 k/s
84.3 k/s	8.2 k/s	896.4 /s	766.7 /s	588.4 Mb/s
568.2 Mb/s	16.7 Mb/s	685.1 kb/s	2.8 Mb/s	

All None Display: Sum Rate

Netflow Processing

Source: peer1 peer2 gateway site upstream
All Sources and <none>

Options:
☐ List Flows ☒ Stat TopN
Top: 10
Stat: Flow Records order by flows
Aggregate: ☒ proto ☒ srcPort ☒ dstPort
Limit: ☐ Packets
Output: line ☐ / IPv6 long
Clear Form process

```
** nfdump -M /netflow0/nfsen-demo/profile-data/live/peer1:peer2:gateway:site:upstream -T -r 2007/05/31/04/nfcapd.200705310440
nfdump filter:
any
Aggregated flows 2797250
Top 10 flows ordered by flows:
Date flow start      Duration Proto      Src IP Addr:Port      Dst IP Addr:Port      Packets  Bytes  Flows
2007-05-31 04:39:54.045 299.034 UDP      116.147.95.88:1110    -> 188.142.64.162:27014    68      5508   68
2007-05-31 04:39:56.282 298.174 UDP      116.147.249.27:1478   -> 188.142.64.163:27014    67      5427   67
2007-05-31 04:39:57.530 298.206 UDP      117.196.44.62:1031    -> 188.142.64.166:27014    67      5427   67
2007-05-31 04:39:57.819 298.112 UDP      117.196.75.134:1146   -> 188.142.64.167:27014    67      5427   67
2007-05-31 04:39:53.787 297.216 UDP      61.191.235.132:4121   -> 60.9.138.37:4121       62      3720   62
2007-05-31 04:39:55.354 300.833 UDP      60.9.138.37:2121     -> 118.25.93.95:2121      61      3660   61
2007-05-31 04:39:58.936 298.977 UDP      60.9.138.36:2121     -> 119.182.123.166:2121   61      3660   61
2007-05-31 04:39:54.329 303.585 UDP      120.150.194.76:2121   -> 60.9.138.37:2121       61      3660   61
2007-05-31 04:39:53.916 300.734 UDP      60.9.138.37:2121     -> 125.167.25.128:2121    61      3660   61
2007-05-31 04:39:57.946 300.353 UDP      60.9.138.36:2121     -> 121.135.4.186:2121     61      3660   61

IP addresses anonymized
Summary: total flows: 4616424, total bytes: 156.6 G, total packets: 172.6 M, avg bps: 644.8 M, avg pps: 90946, avg bpp: 929
Time window: 2007-05-31 04:11:49 - 2007-05-31 04:44:58
Total flows processed: 4616424, skipped: 0, Bytes read: 240064932
Sys: 6.184s flows/second: 746464.4 Wall: 6.185s flows/second: 746361.3
```

nfsen 1.3

Bringing the power of the command line forward

```

My traceroute  [v0.74]
pumba.kramse.dk (::)                               Fri Mar  4 10:22:08 2011
Keys:  Help    Display mode    Restart statistics   Order of fields    quit

          Packets          Pings
Host      Loss%  Snt   Last    Avg    Best   Wrst   StDev
1. 2001:16d8:dd0e:1::100      0.0%    7     0.1     0.1    0.1    0.1    0.0
2. gw-26.cph-01.dk.sixxs.net  0.0%    6    13.7    13.7   13.6   13.9    0.1
3. 3229-sixxs.cr0-r72.gbl-cph.dk.ip6.p80.net  0.0%    6    14.3    14.3   14.3   14.4    0.0
4. te4-3-r72.cr0-r70.tc2-ams.nl.ip6.p80.net  0.0%    6    25.4    51.0   25.3  178.6   62.5
5. 20gigabitethernet1-3.core1.ams1.ipv6.he.net  0.0%    6    25.8    26.5   25.7   29.9    1.7
6. ge-0.ams-ix.amstnl02.nl.bb.gin.ntt.net      0.0%    6    26.3    32.0   26.3   60.2   13.8
7. as-0.r25.tokyjp01.jp.bb.gin.ntt.net         0.0%    6   284.1   306.1  283.6  372.8   37.1
8. po-2.a15.tokyjp01.jp.ra.gin.ntt.net         0.0%    6   298.4   298.3  298.1  298.5    0.2
9. ge-8-2.a15.tokyjp01.jp.ra.gin.ntt.net       0.0%    6   301.2   301.2  300.9  301.7    0.3
10. ve44.foundry6.otemachi.wide.ad.jp          0.0%    6   300.9   300.9  300.8  301.0    0.1
11. ve42.foundry4.nezu.wide.ad.jp              0.0%    6   301.0   301.0  300.9  301.3    0.2
12. cloud-net1.wide.ad.jp                     0.0%    6   301.1   301.0  300.9  301.1    0.1
13. 2001:200:dff:fff1:216:3eff:feb1:44d7      0.0%    6   301.3   301.2  301.0  301.3    0.1
  
```

Nping new kid on the block

```
hlk@pumba:nmap-5.51$ nping www.solidonetworks.com
```

```
Starting Nping 0.5.51 ( http://nmap.org/nping ) at 2011-03-04 10:18 CET
SENT (0.0059s) Starting TCP Handshake > www.solidonetworks.com:80 (91.102.95.20:80)
RECV (0.0067s) Handshake with www.solidonetworks.com:80 (91.102.95.20:80) completed
SENT (1.0093s) Starting TCP Handshake > www.solidonetworks.com:80 (91.102.95.20:80)
RECV (1.0105s) Handshake with www.solidonetworks.com:80 (91.102.95.20:80) completed
SENT (2.0193s) Starting TCP Handshake > www.solidonetworks.com:80 (91.102.95.20:80)
RECV (2.0201s) Handshake with www.solidonetworks.com:80 (91.102.95.20:80) completed
SENT (3.0293s) Starting TCP Handshake > www.solidonetworks.com:80 (91.102.95.20:80)
RECV (3.0302s) Handshake with www.solidonetworks.com:80 (91.102.95.20:80) completed
SENT (4.0393s) Starting TCP Handshake > www.solidonetworks.com:80 (91.102.95.20:80)
RECV (4.0402s) Handshake with www.solidonetworks.com:80 (91.102.95.20:80) completed
```

```
Max rtt: 1.193ms | Min rtt: 0.781ms | Avg rtt: 0.932ms
TCP connection attempts: 5 | Successful connections: 5 | Failed: 0 (0.00%)
Tx time: 4.03457s | Tx bytes/s: 99.14 | Tx pkts/s: 1.24
Rx time: 4.03550s | Rx bytes/s: 49.56 | Rx pkts/s: 1.24
Nping done: 1 IP address pinged in 4.04 seconds
```

Nping is sexy too

```
hlk@pumba:nmap-5.51$ nping -6 www.solidonetworks.com
```

```
Starting Nping 0.5.51 ( http://nmap.org/nping ) at 2011-03-04 10:18 CET
SENT (0.0061s) Starting TCP Handshake > 2a02:9d0:10::9:80
RECV (0.0224s) Handshake with 2a02:9d0:10::9:80 completed
SENT (1.0213s) Starting TCP Handshake > 2a02:9d0:10::9:80
RECV (1.0376s) Handshake with 2a02:9d0:10::9:80 completed
SENT (2.0313s) Starting TCP Handshake > 2a02:9d0:10::9:80
RECV (2.0476s) Handshake with 2a02:9d0:10::9:80 completed
SENT (3.0413s) Starting TCP Handshake > 2a02:9d0:10::9:80
RECV (3.0576s) Handshake with 2a02:9d0:10::9:80 completed
SENT (4.0513s) Starting TCP Handshake > 2a02:9d0:10::9:80
RECV (4.0678s) Handshake with 2a02:9d0:10::9:80 completed
```

```
Max rtt: 16.402ms | Min rtt: 16.249ms | Avg rtt: 16.318ms
TCP connection attempts: 5 | Successful connections: 5 | Failed: 0 (0.00%)
Tx time: 4.04653s | Tx bytes/s: 98.85 | Tx pkts/s: 1.24
Rx time: 4.06292s | Rx bytes/s: 49.23 | Rx pkts/s: 1.23
Nping done: 1 IP address pinged in 4.07 seconds
```

Are you running IPv6?

Please do not buy devices or connections without asking for IPv6!

RANCID - Really Awesome New Cisco conflg Differ

+ Juniper, Dell, ... <http://www.shrubbery.net/rancid/>

Expect, script etc. great for installing devices with common settings

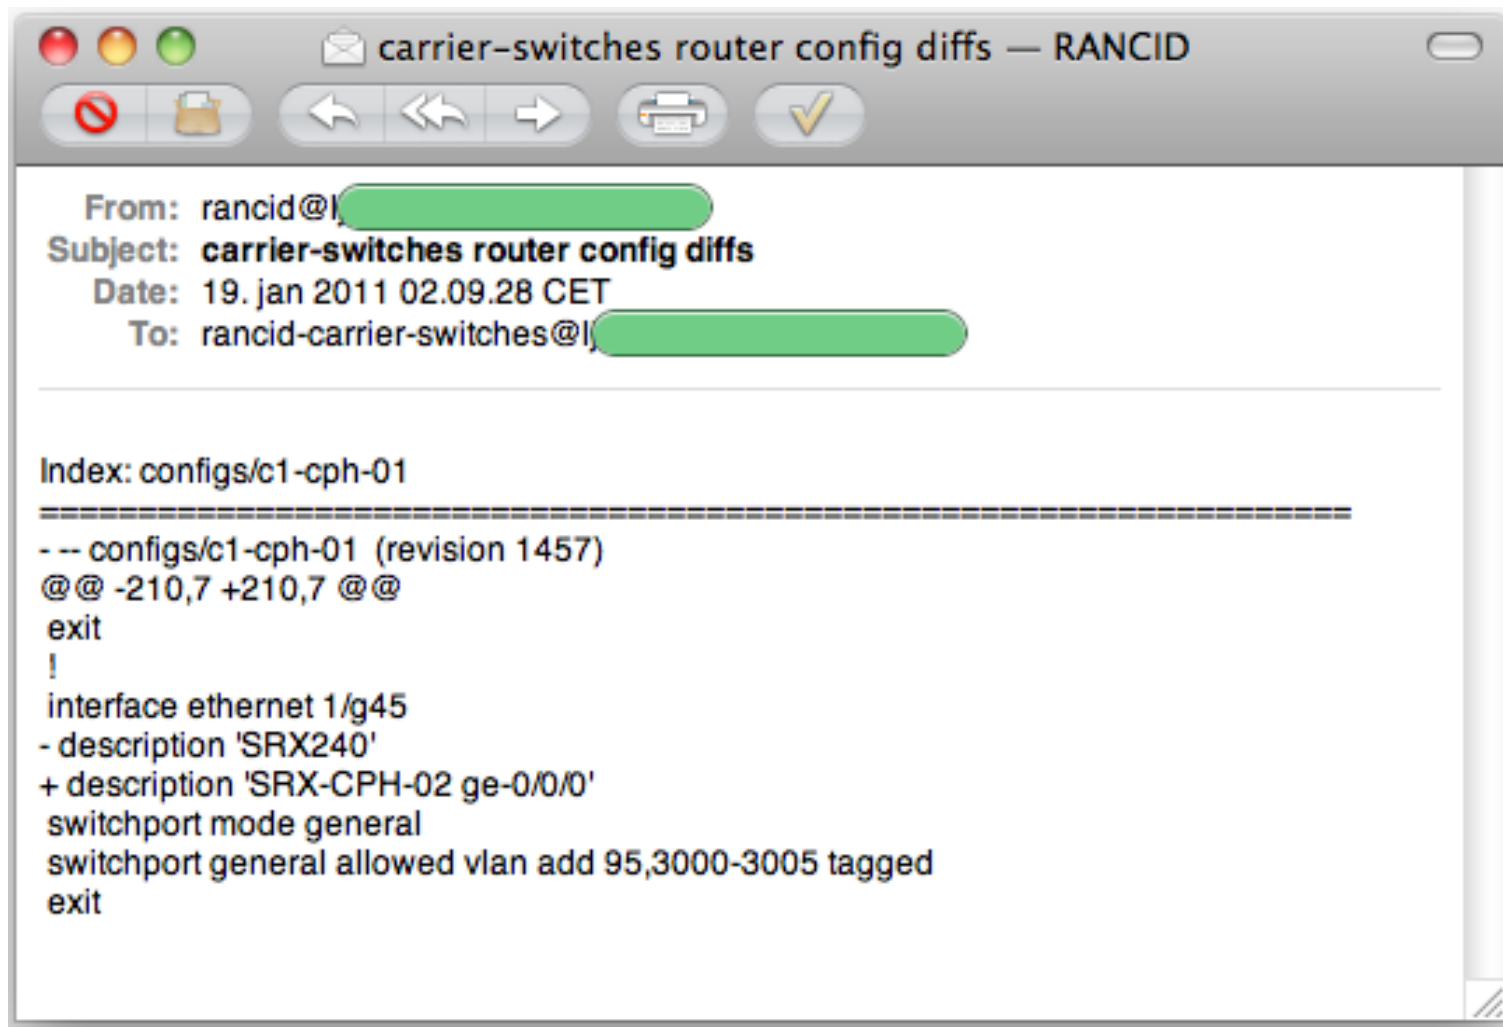
<http://expect.sourceforge.net/> the expect home page

Self discovering: Observium, new and not perfect, but very useful

```
[rancid@ljh routers]$ cat router.db
mx-lux-01:juniper:up
mx-lux-02:juniper:up
...
[rancid@ljh routers]$ crontab -l
# run config differ hourly
07 0-23/2 * * * /usr/local/rancid/bin/rancid-run
# clean out config differ logs
50 23 * * * /usr/bin/find /usr/local/rancid/var/logs -type f -mtime +2 -exec rm {}
```

RANCID will then fetch configurations, and more, and put it into version control SVN/CVS

Changes are emailed to an email alias



SUBVERSION REPOSITORIES RANCID

English - English

(root)/san-switches/configs/san-sw-cph-01 - Rev 1608

Rev

[◀ Rev 1600](#) | [👤 Blame](#) | [📄 Compare with Previous](#) | [🔧 Last modification](#) | [📋 View Log](#) | [📁 Download](#) | [📡 RSS feed](#)

```
!RANCID-CONTENT-TYPE: Dell
!
!
Image Descriptions
image1 : default image
image2 :
Images currently available on Flash
-----
unit      image1      image2      current-active      next-active
-----
```

Hints:

Use rancid user on server and devices, preferably read-only

Use SSH keys to avoid clear text passwords in `~rancid/.cloginrc`

Expose the Subversion to others in the organization using websvn

Me: Why does it take that long to change this setting?

Them: Because we log into each router manually

RANCID uses Expect, for example in the clogin script

Using the clogin script it is possible to perform a command on - say 60 routers in less than 10 minutes ...

Sure, you should watch over the process, but typing your loooong and complex network password 60 times?!

Are you fucking mental?!

Expect example - clogin from RANCID

```
expect {
    -re "(Connection refused|Secure connection \[^\n\r]+ refused)" {
        catch {close}; catch {wait};
        if !$progs {
            send_user "\nError: Connection Refused ($prog): $router\n"
            return 1
        }
    }
    -re "(Connection closed by|Connection to \[^\n\r]+ closed)" {
        catch {close}; catch {wait};
        if !$progs {
            send_user "\nError: Connection closed ($prog): $router\n"
            return 1
        }
    }
    -re "(Host key not found |The authenticity of host .* be established).*\ (yes\|no\)\?" {
        send "yes\r"
        send_user "\nHost $router added to the list of known hosts.\n"
        exp_continue
    }
}
```

Observium is an autodiscovering PHP/MySQL based network monitoring system focused primarily on Cisco and Linux networks but includes support for a wide range of network hardware and operating systems.

Tested it at The Camp summer 2010 not ready

Tested it again Fall 2010, not finished, but useful enough

`http://observium.org/`

Easy up and running

`http://www.observium.org/wiki/CentOS_SVN_Installation`

Why makes Observium great?



```
[root@wiseguy observium]# ./addhost.php
```

Add Host Tool

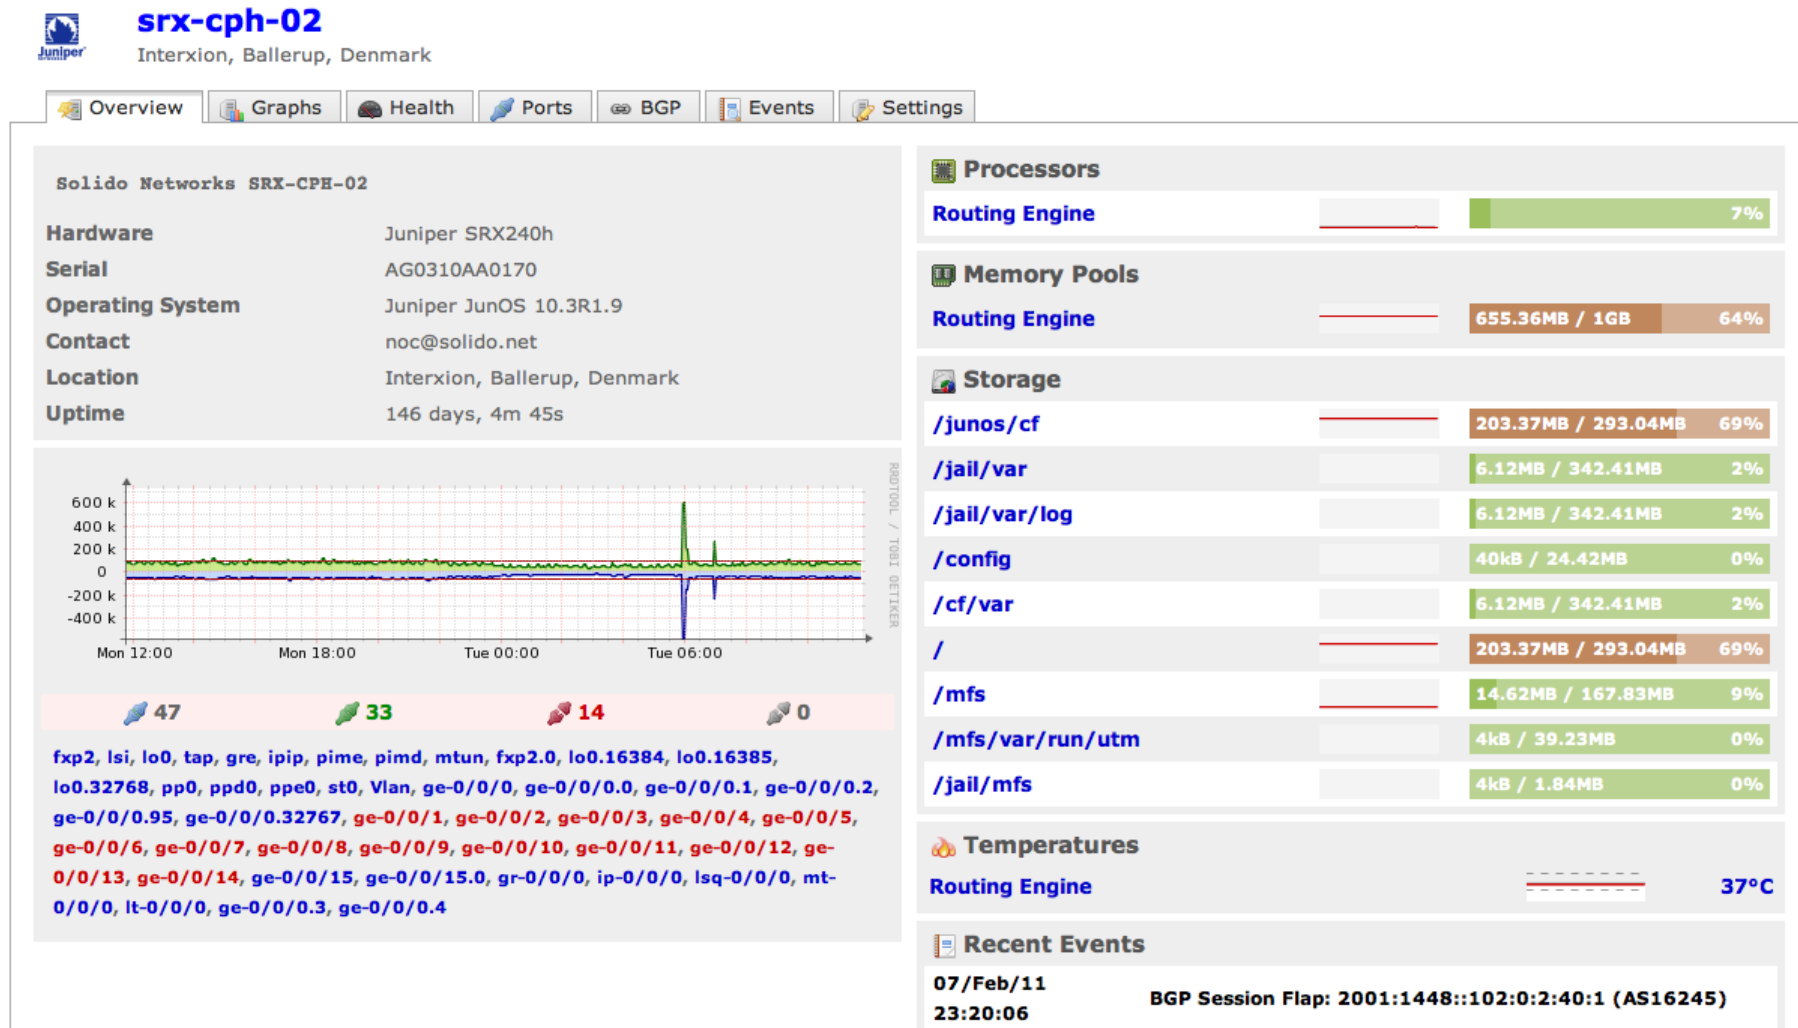
Usage: `./addhost.php <hostname> [community] [v1|v2c] [port]`

Configure your devices in a consistent way

Add host and discover the rest using Observium - done

Surf your network data, including BGP sessions

Observium example router overview



More useful information than default vendor interface! (flash)

```
--- JUNOS 9.6R2.11 built 2009-10-06 20:09:34 UTC  
hlk@ ...> show configuration interfaces
```

```
...  
xe-2/0/0 {  
    description "Transit: Netgroup (AS16245)";  
...  
ge-4/0/0 {  
    description "Cust: xxx (200Mbit contract)";
```

Another router:

```
ge-1/0/1 {  
    unit 0 {  
        description "Peering: LU-CIX Exchange";
```

http://www.observium.org/wiki/Interface_Description_Parsing



Overview
 Devices
 Services
 Locations
 Ports
 Health
 BGP Sessions

srx-cph-03
Device Down
 Interxion, Baller...

Recent Eventlog Entries

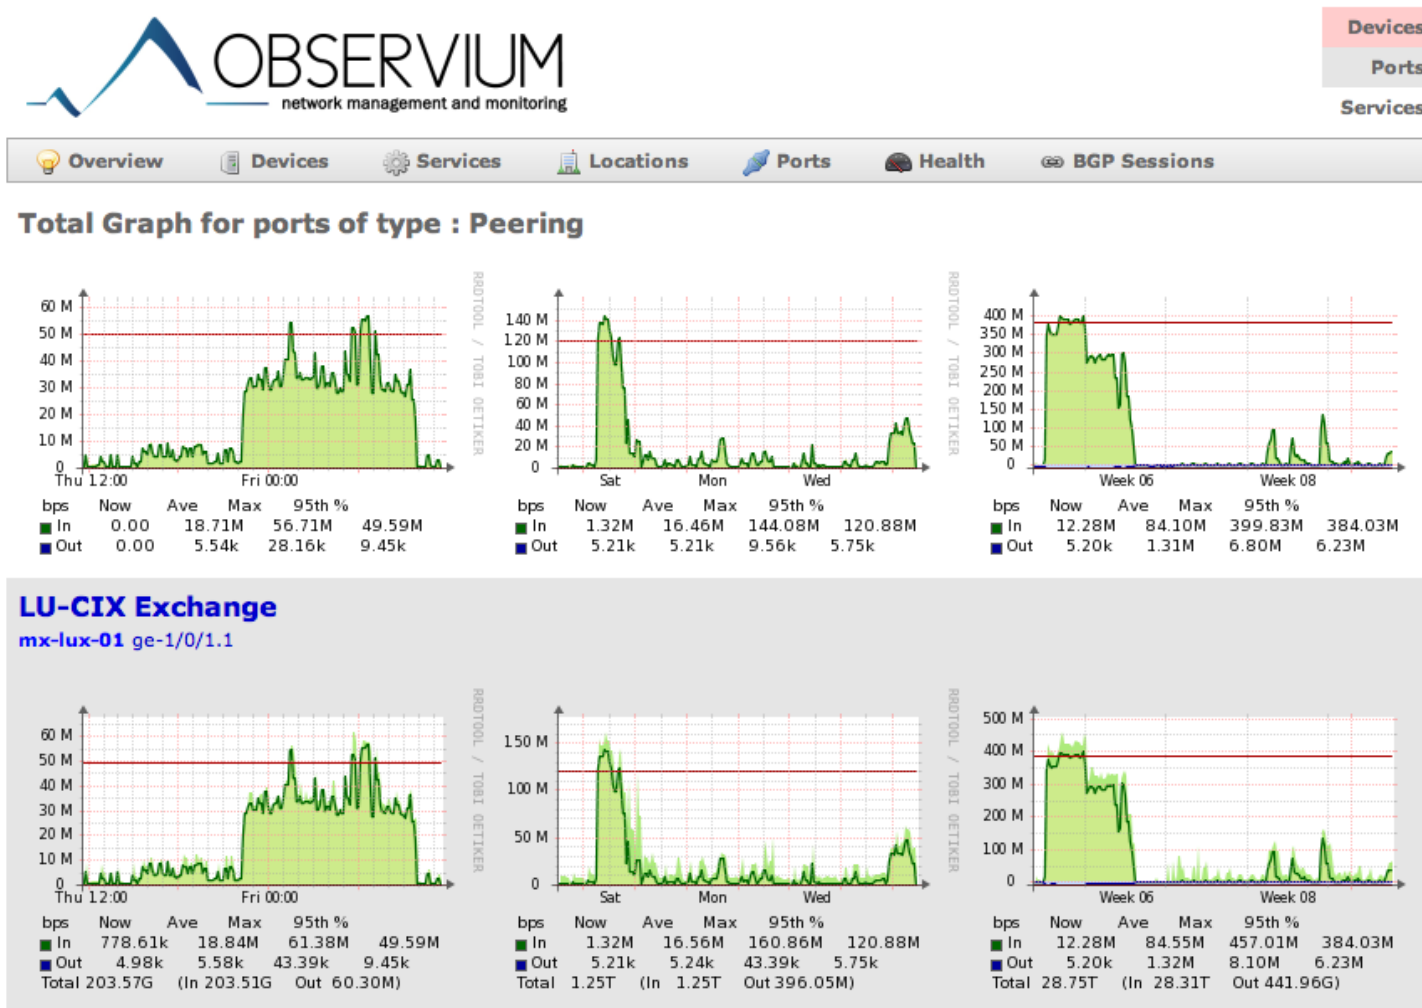
2011-03-04 11:16:02	colo-cph-01	
2011-03-04 11:16:02	colo-cph-01	
2011-03-04 11:16:02	colo-cph-01	
2011-03-04 11:16:02	colo-cph-01	
2011-03-04 08:26:12	blade-sw-cph-02	1/0/4
2011-03-04 08:26:12	blade-sw-cph-02	1/0/4
2011-03-04 08:26:12	blade-sw-cph-02	1/0/4
2011-03-04 08:26:09	blade-sw-cph-01	1/0/4
2011-03-04 08:26:09	blade-sw-cph-01	1/0/4
2011-03-04 08:26:09	blade-sw-cph-01	1/0/4
2011-03-04 08:26:10	blade-sw-cph-02	1/0/4

All Ports
 Errored (3)
 Ignored (1096)
 Pseudowires
 VRFs
 IPv4 Search
 IPv6 Search

Customers
 Transit
Peering
 Peering + Transit
 Core
 Down
 Disabled
 Deleted (499)



1/0/11 -> NULL
 1/0/15 -> NULL
 1/0/19 -> NULL
 1/0/20 -> NULL
 status: down -> up
 0 -> 100000000000
 speed: 0 -> 10000
 status: down -> up
 0 -> 100000000000
 speed: 0 -> 10000
 ifAdminStatus: down -> up

Overview peering ports

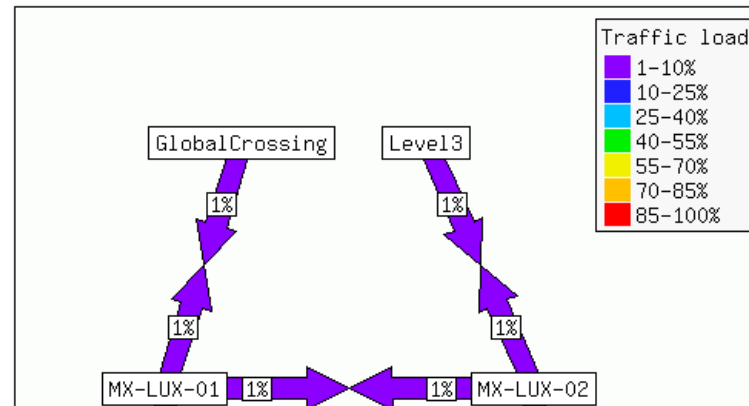


(does not show all peerings we have)

mx-lux-01 :: Port ge-1/0/1.1 :: port_bits



Drill down and hyperlinks everywhere



Reuse data - google: mrtg weathermap will show multiple tools

We use: <http://netmon.grnet.gr/weathermap/>

Many tools use RRDtool - recreate specific graphs

Gather links and create overview pages with the most important ones

A virtual NOC with Open Source IMHO better than any commercial tool

We have shown a number of high quality tools

Use them and keep the flame burning

Open Source is critical and we need more network skills

Dont forget about the nice websites that work to your advantage:

Routing Information Service and Whois

<http://www.ripe.net>

traceroutes to your network

<http://www.traceroute.org>

http://nanog.cluepon.net/index.php/Tools_and_Resources

<http://labs.ripe.net/Members/vastur/the-shape-of-a-bgp-update>

<http://www.delicious.com/kramshoej/netflow> and other tags

And use google :-)

Henrik Lund Kramshøj
hlk@solidonetworks.com

`http://www.solidonetworks.com`

You are always welcome to send me questions later via email

Networks tools are here already - use them



- Henrik Lund Kramshøj, IT-security and IP network consultant
- Email: hlik@solidonetworks.com Mobile: +45 2026 6000
- Educated from the Computer Science Department at the University of Copenhagen, DIKU
- CISSP and CEH certified
- 2003 - 2010 Independent security consultant
- 2010 - owner and partner in Solido Networks ApS